

Adnan Mashrur Sadad

AI / ML ENGINEER | APPLIED AI, EVALUATION AND ML SYSTEMS

Kuala Lumpur, Malaysia | +60 11-3968 7435 | adnanmashrursadad@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

AI/ML engineer with enterprise development experience at Joget and first-author RAG and Text-to-SQL research. Builds evaluated AI applications and deployable ML pipelines.

TECHNICAL SKILLS

AI and ML: Python, SQL, scikit-learn, XGBoost, pandas, NumPy, Hugging Face, RAG, LLM evaluation

Applications: FastAPI, Pydantic, SQLAlchemy, React, TypeScript, Groq, structured-output validation

ML Infrastructure: Docker, Kubernetes, Kafka/Redpanda, PyFlink, Feast, MLflow, ONNX, GitHub Actions

EXPERIENCE

Joget Inc. | Software Developer Intern

Oct 2024 - Mar 2025

- Built and deployed Joget DX workflow automations and UI components for client teams; collaborated with senior engineers, QA and product managers in Agile sprints.
- Integrated REST APIs to remove manual re-entry steps and maintained modules through recurring production release cycles.

SELECTED PROJECTS

Driftline | Streaming Fraud Detection and MLOps | [GitHub](#)

2026

- Replayed 590K IEEE-CIS transactions through Redpanda/PyFlink into Redis/Parquet stores; measured XGBoost PR-AUC of 0.4751 on a chronological split versus 0.6974 on a random split.
- Implemented PSI-triggered retraining with a shadow promotion gate: candidate PR-AUC 0.4802 versus incumbent 0.2811 on the same holdout; both triggered retrains passed.
- Served ordinal-encoded XGBoost via ONNX/FastAPI: 10,018 requests with zero failures and 380 ms p95 at 50 concurrent users on a 4-vCPU GCP VM.

InterviewPilot | Applied AI Interview Platform | [GitHub](#)

2026

- Built a React/FastAPI application with persistent interviews, Groq transcription, schema-validated answer evaluation and repair loops; bounded follow-ups and disclosed fallback scoring when AI is unavailable.
- Verified 19 backend tests with 89% statement coverage and CI browser checks from session creation to report generation; external AI calls mocked or disabled in tests.

Freight Rate ML | Temporal Regression | [GitHub](#)

2026

- Benchmarked linear and boosted models on a 9,523-row temporal holdout with training-only preprocessing: boosted MAPE 6.89% versus 8.28%; linear baseline retained better MAE/RMSE.

RESEARCH

SAFE-RAG | Regulatory RAG Evaluation | [GitHub](#)

First author | Short paper under review at ALTA 2026

- Estimated an 8.6% wrong-scope grounding rate among schema-valid, citation-entailed answers (95% CI: 4.5-13.5%; 200 annotated items); revised-round kappa: 0.786 (n=30).

Text-to-SQL Answerability under Schema Evolution

First author | Manuscript complete; not yet submitted

- Built a 1,382-instance benchmark across 20 Spider databases and seven drift families; ran 27,384 generations over 39 experimental cells on Kaggle T4 GPUs.
- Detected 79.5% of structurally invisible drift at a 13.0% false-alarm rate, or 51.2% at zero false alarms; evaluated against 5,627 benign updates (annotation kappa: 0.918).

EDUCATION

Universiti Teknologi Malaysia | MJIT

2021 - 2026

B.Sc. Computer Science (Software Engineering) with Honours | CGPA: 3.43/4.00

Dean's List: 2022/2023, 2025/2026 | Gold Medal and Best Video Award, myHCI-UX 2025