

Adnan Mashrur Sadad

AI / ML ENGINEER | APPLIED AI, EVALUATION AND ML SYSTEMS

Kuala Lumpur, Malaysia | +60 11-3968 7435 | adnanmashrursadad@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

AI/ML engineer and first-author researcher building evaluated LLM systems and deployable ML pipelines, with enterprise software experience at Joget.

TECHNICAL SKILLS

AI and ML: Python, SQL, scikit-learn, XGBoost, PyTorch, pandas, NumPy, Hugging Face, RAG, Text-to-SQL, LLM evaluation

Applications: FastAPI, Pydantic, SQLAlchemy, React, TypeScript, Groq, agentic workflows, structured-output validation

ML Infrastructure: Docker, Kubernetes, Kafka/Redpanda, PyFlink, Feast, MLflow, ONNX, GitHub Actions

EXPERIENCE

Software Developer Intern

Oct. 2024 – Mar. 2025

Joget Inc.

Kuala Lumpur, Malaysia

- Built and deployed Joget DX workflow automations and UI components for client teams; collaborated with senior engineers, QA, and product managers in Agile sprints.
- Integrated REST APIs to remove manual re-entry steps and maintained modules through recurring production release cycles.

SELECTED PROJECTS

Driftline | *Streaming Fraud Detection and MLOps* | [GitHub](#)

2026

- Replayed 590K IEEE-CIS transactions through Redpanda/PyFlink; engineered a PyTorch Geometric GraphSAGE model over a 605K-node identity graph with a unit-tested leakage boundary; measured XGBoost PR-AUC of 0.4751 on a chronological split versus 0.6974 on a random split.
- Implemented PSI-triggered retraining with a shadow promotion gate: candidate PR-AUC 0.4802 versus incumbent 0.2811 on the same holdout; both triggered retrains passed.
- Served ordinal-encoded XGBoost via ONNX/FastAPI: 10,018 requests with zero failures and 380 ms p95 at 50 concurrent users on a 4-vCPU GCP VM.

InterviewPilot | *Applied AI Interview Platform* | [GitHub](#)

2026

- Built a React/FastAPI application with persistent interviews, Groq transcription, schema-validated answer evaluation and repair loops; bounded agentic follow-ups and disclosed fallback scoring when AI is unavailable.
- Verified 19 backend tests with 89% statement coverage and CI browser checks from session creation to report generation; external AI calls mocked or disabled in tests.

Freight Rate ML | *Temporal Regression* | [GitHub](#)

2026

- Benchmarked linear and boosted models on a 9,523-row temporal holdout with training-only preprocessing: boosted MAPE 6.89% versus 8.28%; linear baseline retained better MAE/RMSE.

RESEARCH

SAFE-RAG | *Deceptive Grounding in Regulatory RAG* | [GitHub](#)

2025 – 2026

- First author; short paper under review at ALTA 2026. Estimated an 8.6% wrong-scope grounding rate among schema-valid, citation-entailed answers (95% CI: 4.5–13.5%; 200 annotated items); revised-round κ : 0.786 ($n = 30$).

Text-to-SQL Answerability under Schema Evolution

2026

- First author; manuscript complete, not yet submitted. Built a 1,382-instance benchmark across 20 Spider databases and seven drift families; ran 27,384 generations over 39 experimental cells on Kaggle T4 GPUs.
- Detected 79.5% of structurally invisible drift at a 13.0% false-alarm rate, or 51.2% at zero false alarms; evaluated against 5,627 benign updates (annotation κ : 0.918).

EDUCATION

Universiti Teknologi Malaysia (UTM), MJIIT

2021 – 2026

B.Sc. Computer Science (Software Engineering) with Honours | CGPA: 3.43/4.00

Kuala Lumpur, Malaysia

- Dean's List: 2022/2023, 2025/2026.
- Thesis: OCR-based personal finance tracking for automated income-tax-ready expense classification.

AWARDS & LEADERSHIP

Gold Medal & Best Video Award, myHCI-UX Student Design Challenge (Oct 2025) — national inter-university UX design competition.

Secretary Treasurer, SOFEA Society (2022–2023); organised AI workshops for 200+ students.